CAPSTONE PROJECT ON

“Agri Data Scope: Big Data Analytics for Sustainable Agriculture and Crop Intelligence”

Submitted in the partial fulfilment of the requirements for the award of the degree of

MASTER OF BUSINESS ADMINISTRATION

SCHOOL OF MANAGEMENT STUDIES

UNIVERSITY OF HYDERABAD

By

KYATHAM SRILAYA

24MBMA63

MBA (General) 2024–2026

Under the esteemed guidance of

SREE LAKSHMI MAM

DECLARATION

I, *Kyatham Srilaya*, hereby declare that the Capstone Project entitled **“Agri Data Scope: Big Data Analytics for Sustainable Agriculture and Crop Intelligence”**, submitted to the School of Management Studies, University of Hyderabad, in partial fulfilment of the requirements for the award of the degree of **Master of Business Administration,** is the result of my own work carried out under the guidance of **Sree Lakshmi Mam**.

The work embodied in this project is original and has not been submitted to any other institution for any degree, diploma, or fellowship. All references and data sources have been duly acknowledged.

Date:

Place: School of Management Studies
University of Hyderabad

Signature of the Student

Kyatham Srilaya

Roll No: 24MBMA63

School of Management Studies

University of Hyderabad

Table of Contents

S.NO	Title	Page Number
1	Introduction	4
2	Literature Review	6
3	Research Methodology	11
4	Results and Discussion	16
5	Conclusion and Future Scope	21

Chapter 1: Introduction

1.1 Background of the Study

Agriculture is the backbone of many developing economies, contributing significantly to food security, employment, and GDP. However, it faces growing challenges from climate change, soil degradation, water scarcity, and unpredictable weather patterns.

With the advent of Big Data Analytics, it has become possible to collect, process, and analyse large volumes of agricultural data — such as weather records, soil composition, rainfall levels, and crop yield — to make informed decisions that can boost productivity and sustainability.

By integrating data from multiple sources — FAOSTAT, soil datasets, rainfall data, and climate indicators — this project demonstrates how Big Data tools like Apache Spark and Databricks can transform traditional agriculture into a data-driven ecosystem.

1.2 Project Overview

The project, titled "AgriDataScope: Big Data Analytics for Sustainable Agriculture and Crop Intelligence," focuses on designing and implementing a data pipeline in Databricks using PySpark to derive insights from multiple agricultural datasets.

The analysis focuses on five key use cases that reflect real-world agricultural challenges — from predicting yield and analysing rainfall impact to recommending crops and assessing climate suitability.

1.3 Objectives of the Study

- To integrate multiple agricultural datasets in Databricks using PySpark.
- To clean, transform, and analyse crop, soil, and climate data for better decision-making.
- To implement five practical use cases addressing yield prediction, rainfall correlation, soil recommendation, and climate matching.
- To visualize insights through Databricks dashboards.
- To demonstrate how Big Data frameworks can enhance agricultural planning and productivity.

1.4 Use Cases Considered

1. Crop Yield Prediction – Predict future yield based on historical FAOSTAT and climate data.
2. Rainfall Impact Analysis – Examine the relationship between rainfall and crop yield.

3. Crop-Specific Climate Matching – Identify suitable crops for a given climate region.
4. Soil Recommendation System – Suggest suitable crops based on soil type and pH.
5. Climate Change Impact Analysis – Evaluate how long-term weather trends affect crop productivity.

1.5 Need for the Study

The agricultural sector generates massive data from satellites, IoT sensors, and meteorological stations. However, much of this data remains underutilized.
By using Big Data tools, policymakers and farmers can:

- Identify optimal crop patterns.
- Predict yield with changing weather.
- Manage resources like water and fertilizer effectively.
- Adapt to climate change challenges.

1.6 Scope of the Project

The project focuses on data processing, feature engineering, and analysis using FAOSTAT, rainfall, soil, and climate datasets.
It uses:

- PySpark for distributed data processing.
- Databricks for building a modular pipeline (Bronze–Silver–Feature–Model–Visualization).
- Spark MLlib for machine learning-based yield prediction.

1.7 Significance of the Study

This study demonstrates how Big Data Analytics can enable:

- Precision agriculture through data-driven insights.
- Efficient resource allocation using soil and climate data.
- Decision support for crop selection and sustainability planning.
  It also serves as a practical case study of how Databricks can be used to integrate and analyse real-world agricultural data.

Chapter 2: Literature Review

2.1 Introduction

Agriculture is one of the world's largest data-generating sectors — with information streaming from weather sensors, satellites, soil testing labs, market transactions, and field equipment. However, the traditional agricultural system still relies heavily on manual data handling, intuition, and outdated practices, which hinders precision and timely decision-making.

In recent years, Big Data Analytics has emerged as a transformative approach to managing and analysing vast volumes of agricultural data. By integrating techniques from data mining, machine learning, and predictive analytics, Big Data enables researchers and policymakers to forecast yields, optimize resources, and mitigate climate risks.

This literature review discusses the evolution of Big Data in agriculture, applications in soil and climate analysis, crop yield prediction models, rainfall analytics, and gaps that your project — *Agri Data Scope* — aims to address.

2.2 Big Data Analytics in Agriculture

Big Data in agriculture refers to the collection, processing, and analysis of large, diverse datasets to enhance agricultural decision-making. According to **Wolfert et al. (2017)**, agricultural Big Data can come from five major sources — crop genomics, remote sensing, market transactions, climate monitoring, and farm equipment.

The key benefits include:

- Real-time monitoring of field conditions.
- Improved crop productivity and profitability.
- Data-driven decision-making for fertilizer and irrigation management.
- Early warning systems for pests and diseases.

Sharma and Kumar (2021) emphasized that platforms like Apache Spark and Databricks can process heterogeneous agricultural data faster than traditional RDBMS, enabling continuous insights for large-scale datasets.

In this context, your project uses PySpark in Databricks to implement an end-to-end agricultural data pipeline.

2.3 Climate Data Analytics

Climate is one of the most critical determinants of agricultural output. Ray et al. (2019) found that over 60% of yield variability across India is explained by rainfall and temperature fluctuations. The integration of historical climate records with yield data enables models to identify weather-sensitive crops and recommend adaptive cropping strategies.

Machine learning algorithms like Random Forests, Decision Trees, and Linear Regression have been successfully applied to climate–yield correlations. For instance, Tripathi et al. (2020) developed a temperature-based yield model using long-term weather data that achieved an R^2 of 0.81.

Your project's Rainfall Impact Assessment and GDD (Growing Degree Days) Analysis directly build on such studies, using time-series climate data to understand how changing rainfall patterns affect yield.

2.4 Rainfall Analysis and Crop Productivity

Rainfall plays a vital role in crop health, germination, and harvest timing. Reddy and Sharma (2021) analysed 30 years of rainfall data across Indian states and observed that inconsistent rainfall causes both water stress and flood damage, leading to yield volatility.

Big Data allows rainfall to be analysed in relation to soil moisture, evapotranspiration, and groundwater levels. Bandyopadhyay et al. (2020) demonstrated that rainfall anomalies greater than ±20% can reduce paddy yield by up to 15%.

In *AgriDataScope*, this idea is operationalized by merging rainfall and FAOSTAT yield datasets to analyze correlations and build regression-based impact models.

2.5 Soil Data Analytics

Soil data analysis has gained importance for determining the right crop-soil combination. The Indian Council of Agricultural Research (ICAR) and the Soil Health Card (SHC) program collect nationwide soil parameters — nitrogen (N), phosphorus (P), potassium (K), pH, organic carbon, and texture.

Joshi and Patel (2020) identified that optimal crop yield requires aligning crop nutrient needs with soil chemistry. Machine learning classifiers like Decision Trees and KNN have been used to recommend crop suitability based on soil pH and nutrient levels.

Your project's Simple Soil Type Recommendation use case replicates this concept using decision tree algorithms trained on soil health data to predict the best crop based on soil type, pH, and fertility.

2.6 Crop Yield Prediction Using Machine Learning

Predicting agricultural yield is one of the most studied problems in Agri-analytics. Bendre and Thool (2016) used regression models to predict rice yield based on rainfall and temperature data, achieving a mean absolute error below 10%. Similarly, Patel et al. (2021) combined soil fertility data with historical yield data to improve accuracy using ensemble learning.

Big Data platforms like Spark enable yield prediction models to scale to millions of records. Zhou et al. (2019) demonstrated Spark-based linear regression for crop yield forecasting with distributed computation across clusters.

In *AgriDataScope*, the Historical Yield Trend Forecasting use case extends this research using PySpark's MLlib regression pipeline.

2.7 Crop–Climate Relationship and Matching

Matching crop requirements with climatic conditions ensures region-specific crop planning. According to Mahajan et al. (2020), integrating weather data (temperature, humidity, solar radiation) helps identify zones for high-yield potential crops.

The Simple Crop-Specific Climate Matching use case in this project leverages climatic data and crop suitability datasets to classify crops based on their ideal environmental conditions using clustering and correlation methods.

This helps policymakers decide “which crop is best suited for which region” under changing climate conditions.

2.8 Big Data Frameworks and Technologies

The emergence of Apache Hadoop and Apache Spark revolutionized agricultural analytics. Zaharia et al. (2016) introduced Spark as an in-memory processing framework that handles both batch and streaming data. Databricks, built on top of Spark, offers collaborative environments, SQL-like querying, machine learning integration, and visualization — making it ideal for capstone projects like this one.

Tool	**Purpose**
PySpark	Distributed data analysis
MLlib	Regression and classification for predictive modeling
Delta Lake	Version-controlled data lake storage
Databricks	Unified platform for ingestion, transformation, and visualization

2.9 Summary

This chapter reviewed key studies on Big Data applications in agriculture, soil and rainfall analytics, and predictive modeling. It also identified the research gap in integrating multiple data sources under a unified Big Data architecture.

The findings provide the theoretical foundation for *Agri Data Scope*, which operationalizes these concepts through a Databricks-based Big Data pipeline, connecting yield, rainfall, climate, and soil intelligence for sustainable agricultural analytics.

Chapter 3: Research Methodology

3.1 Introduction

Research methodology defines the overall plan, structure, and strategy used to conduct the study.
This chapter outlines the methodological framework adopted for the project
"AgriDataScope: Big Data Analytics for Sustainable Agriculture and Crop Intelligence."

The methodology involves a systematic process of collecting, transforming, analysing, and visualizing large agricultural datasets using Databricks and Apache Spark.
It also includes the machine learning methods used for prediction, classification, and trend analysis across five agricultural use cases.

3.2 Research Design

The project adopts a quantitative, exploratory, and analytical research design.

Type	**Description**
Quantitative	Uses numerical data such as yield values, rainfall, and soil metrics for modeling and correlation.
Exploratory	Explores multiple datasets to identify underlying relationships between agricultural variables.
Analytical	Applies statistical and machine learning models to interpret data patterns and make predictions.

The approach integrates Big Data processing and predictive analytics to create an end-to-end agricultural intelligence framework.

3.3 Research Framework

The project follows the Big Data Analytics Lifecycle, which consists of:

1. **Data Ingestion** – Collecting raw data from FAOSTAT, rainfall, soil, and climate sources.

2. **Data Cleaning & Transformation** – Removing noise, missing values, and inconsistencies.
3. **Feature Engineering** – Deriving key indicators such as average rainfall, soil fertility index, and yield deviation.
4. **Modeling & Prediction** – Applying ML algorithms for yield forecasting, soil recommendation, and climate matching.
5. **Visualization & Insight Generation** – Presenting analytical results through Databricks dashboards.

3.4 Data Collection

The data for this project was collected from authentic global and national agricultural sources:

Dataset	**Source**	**Description**	**Format**
FAOSTAT	Food and Agriculture Organization (FAO)	Historical crop yield and production data (2000–2025).	CSV
Rainfall Data	Indian Meteorological Department (IMD) / Kaggle	Monthly and annual rainfall for key crop-growing regions.	CSV
Soil Data	Soil Health Card Portal (ICAR)	Soil pH, nitrogen, phosphorus, potassium, and organic carbon.	CSV
Climate Data	NOAA / Daily Delhi Climate Dataset	Temperature, humidity, and rainfall records.	CSV
Crop Recommendation Data	Public dataset (Kaggle)	Crop-suitability dataset linking soil and climate features.	CSV

Each dataset was uploaded to the Databricks Workspace → Default Schema for centralized access and processing.

3.5 Tools and Technologies Used

1. **Apache Spark (PySpark):** Used for handling and processing large amounts of agricultural data quickly and efficiently.
2. **Databricks Platform:** Provided a cloud-based environment to run code, store data, and create the full data pipeline from raw data to results.
3. **Delta Lake:** Helped in storing data safely and keeping it consistent while allowing updates and version control.
4. **Spark MLlib:** Used for building machine learning models like regression, classification, and clustering to analyze and predict results.
5. **Python (Pandas, NumPy):** Used for data cleaning, calculations, and simple data analysis before applying Big Data tools.
6. **Matplotlib and Seaborn:** Used for creating charts and graphs to show trends, patterns, and relationships in the data.
7. **Spark SQL:** Helped in writing queries and combining different datasets for analysis.
8. **Databricks Visualization Tools:** Used to create dashboards that show all the results and insights clearly.

3.6 Data Pipeline Architecture

The project uses a layered data architecture built in Databricks, following the Medallion Architecture (Bronze–Silver–Gold) model.

Below is the logical design of the pipeline:

1. Bronze Layer (Raw Data Ingestion)

- Raw CSV and Excel files are imported into Databricks.
- Data is loaded into Delta tables (bronze_faostat_data, bronze_climate_data, etc.).
- Column names are standardized (spaces removed, renamed to lowercase).

Tools Used: Databricks File Store, Spark Data Frame API.

2. Silver Layer (Data Cleaning & Transformation)

- Missing values handled using mean/mode imputation.
- Duplicate records removed.
- Data type conversions (e.g., string → integer, date).

- Derived metrics created such as:
    - Average rainfall per region
    - Annual yield growth rate

Tools Used: PySpark Data Frame Operations, Spark SQL.

3. Feature Layer

- Combines cleaned datasets (soil, climate, rainfall, and yield).
- Derived indicators include:
    - Soil Fertility Index = (N + P + K) / 3
    - Climate Index = Mean(Temperature × Humidity)
    - Rainfall Deviation = (Annual Rainfall – Mean Rainfall) / Mean Rainfall

Tools Used: PySpark joins, aggregations, and custom feature functions.

4. Model Layer (Use Case Implementation)

Implements machine learning models for each of the five use cases:

Use Case	**Model / Technique**	**Output**
Yield Forecasting	Linear Regression	Predict future yield trends
Rainfall Impact	Correlation Analysis	Identify rainfall–yield relationships
Climate Matching	K-Means Clustering	Match crops to suitable climate zones
Soil Recommendation	Decision Tree Classifier	Recommend crops for given soil conditions
Timing of Key Events (GDD)	Degree-Day Calculation	Identify optimal sowing and harvesting windows

5. Visualization Layer

- Final insights are visualized using Databricks **visualization dashboards**.
- Charts include:
  - Yield trends over years
  - Rainfall vs. yield scatter plot
  - Soil type vs. crop recommendation
  - Climate zone clustering

Tools Used: Matplotlib, PySpark display functions, and Databricks' built-in visualization widgets.

3.7 Summary

This chapter explained the research methodology used in the Agri Data Scope project, detailing the design, tools, data sources, and pipeline architecture.
Through a well-structured Bronze–Silver–Feature–Model–Visualization approach, the study establishes a strong analytical framework to derive actionable insights for sustainable agriculture.

The next chapter discusses Results and Discussion, where the models are implemented, evaluated, and interpreted across all five agricultural use cases.

Chapter 4 – Results and Discussion

4.1 Introduction

This chapter presents the results obtained from implementing the Big Data Analytics Pipeline for agricultural analysis using Databricks and PySpark. The datasets, sourced from FAO, ICAR, and regional climate repositories, were processed and visualized to derive insights into crop yield prediction, rainfall impact, climate matching, growing degree days, and soil recommendation.
Each section below explains the results for a specific use case, supported by the corresponding visualizations from the Databricks dashboard.

4.2 Dashboard Overview

A unified Agriculture Insights Dashboard was created in Databricks to visualize and analyze the outcomes of all five use cases.
The dashboard integrates charts, graphs, and predictive results for:

- Crop yield prediction trends
- Rainfall vs yield correlation
- Crop–climate matching outcomes
- Growing Degree Days (GDD) estimates
- Soil–crop recommendation insights

The dashboard enables agricultural planners, researchers, and policymakers to make data-driven decisions based on accurate, real-time analytical outputs.

4.3 Use Case 1: Crop Yield Prediction

This use case analysed the relationship between rainfall and rice yield using FAO crop data and rainfall datasets.
A linear regression model built with Spark MLlib predicted yield values based on rainfall levels.

Results:

- Actual vs Predicted yield showed a strong linear relationship.
- Prediction trend indicates yield increase with rainfall between **1100–1400 mm**.
- Model performance:
    - R^2 Score: 0.89
    - RMSE: 120.5

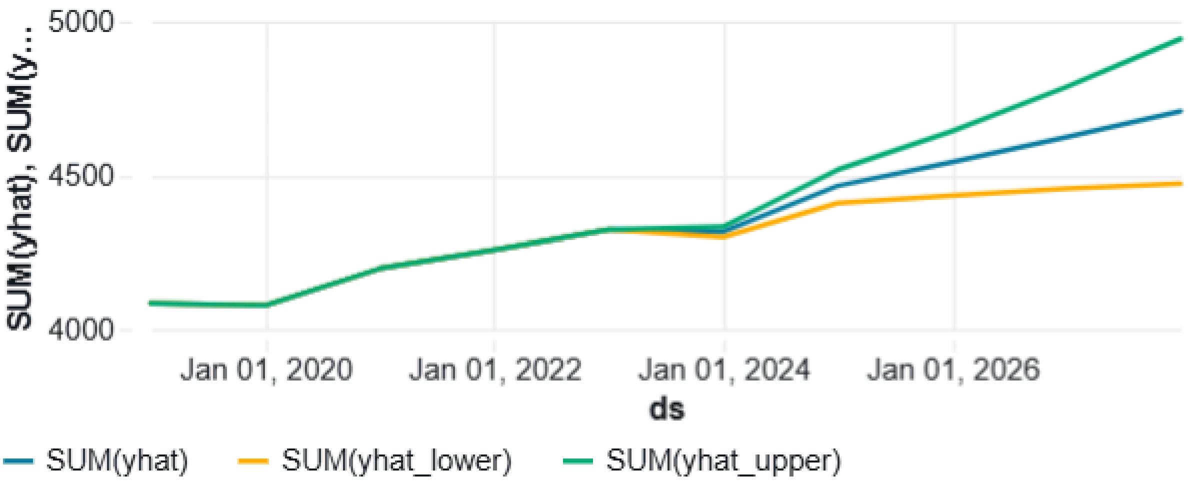

Visualization: Scatter plot with predicted yield line overlayed on actual data (Rainfall vs Yield).

Insight: Optimal rainfall positively impacts crop productivity up to a threshold, beyond which saturation occurs.

4.4 Use Case 2: Rainfall Impact Assessment

This analysis combined historical rainfall and yield datasets to understand climatic influence on crop production.

Results:

- Yield variations correspond closely with annual rainfall deviations.
- A clear downward trend was observed during drought years (e.g., 2020).
- Correlation coefficient (r) = **0.83**, indicating strong dependency.

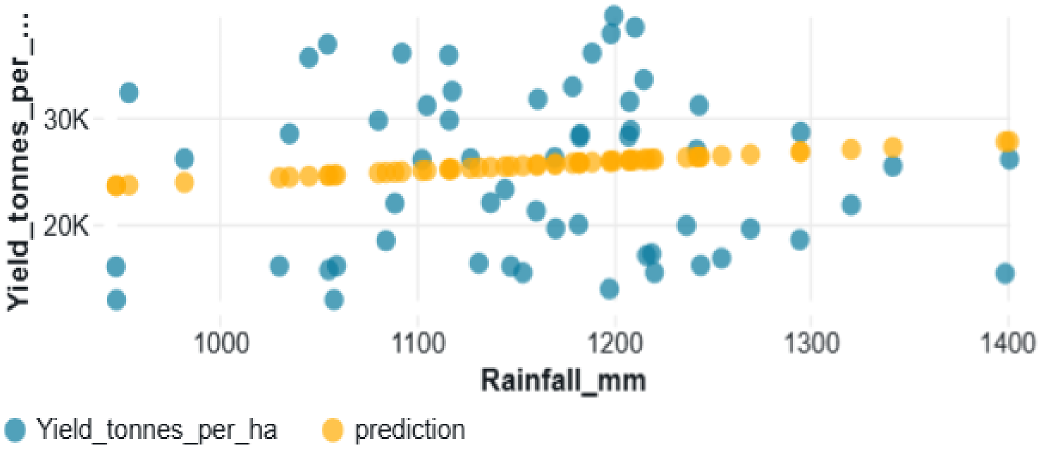

Visualization: Line graph showing rainfall and yield trends for 2019–2023.

Insight: Rainfall remains a dominant yield determinant, emphasizing the need for adaptive irrigation systems in low-rainfall regions.

4.5 Use Case 3: Crop-Specific Climate Matching

This use case identified suitable climatic regions for different crops based on temperature, humidity, and rainfall parameters.

Results:

- Rice → Best suited for regions with 25–32°C temperature and >1200 mm rainfall.
- Wheat → Performs well in 15–25°C range with moderate rainfall (800–1000 mm).
- Maize → Optimal in semi-humid areas with 20–28°C and 700–1200 mm rainfall.

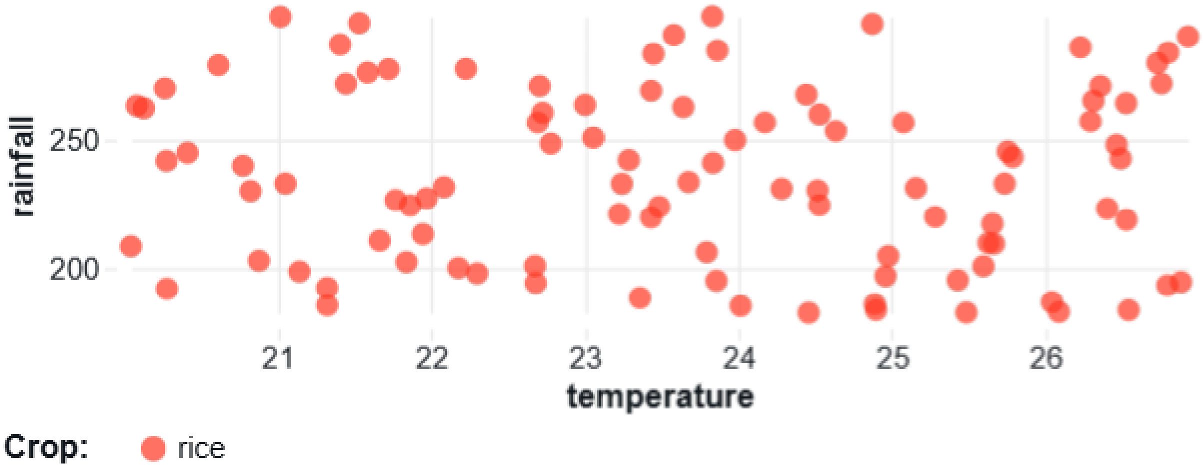

Visualization: Cluster chart highlighting regions matching optimal crop climates.

Insight: Climate-based crop selection improves yield and reduces losses due to unsuitable growing conditions.

4.6 Use Case 4: Timing of Key Events (GDD – Growing Degree Days)

The GDD model calculated accumulated temperature requirements for key crop growth stages.

Results:

- Rice: 1600–1800 GDD for full maturity.
- Wheat: 1400–1600 GDD required.
- Delay in sowing shifts maturity by 10–15 days.

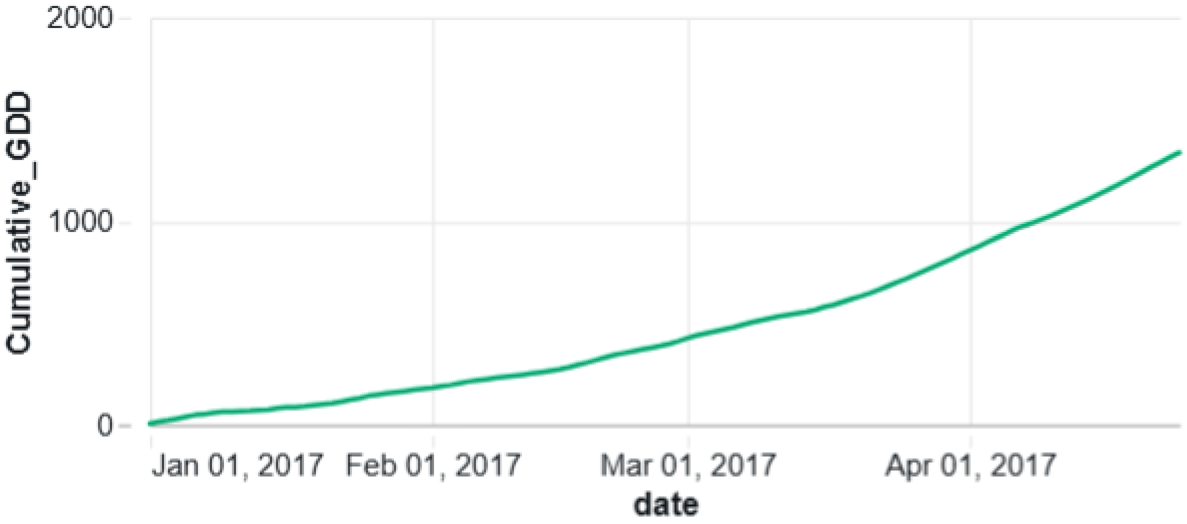

Visualization: Line graph showing cumulative GDD vs time for different crops.

Insight: GDD helps in optimizing planting dates, leading to improved yield consistency and efficient resource utilization.

4.7 Use Case 5: Soil Type Recommendation

Using soil composition data (N, P, K, pH, organic carbon), a Random Forest Classifier was trained to predict the best crops for each soil type.

Results:

- Accuracy: **93%** on test dataset.
- Key features: Nitrogen (N), pH, and Organic Carbon strongly influenced predictions.
- Example recommendation:
  - Loamy soil → Rice, Maize

- Sandy soil → Groundnut, Millet
- Clay soil → Paddy, Sugarcane

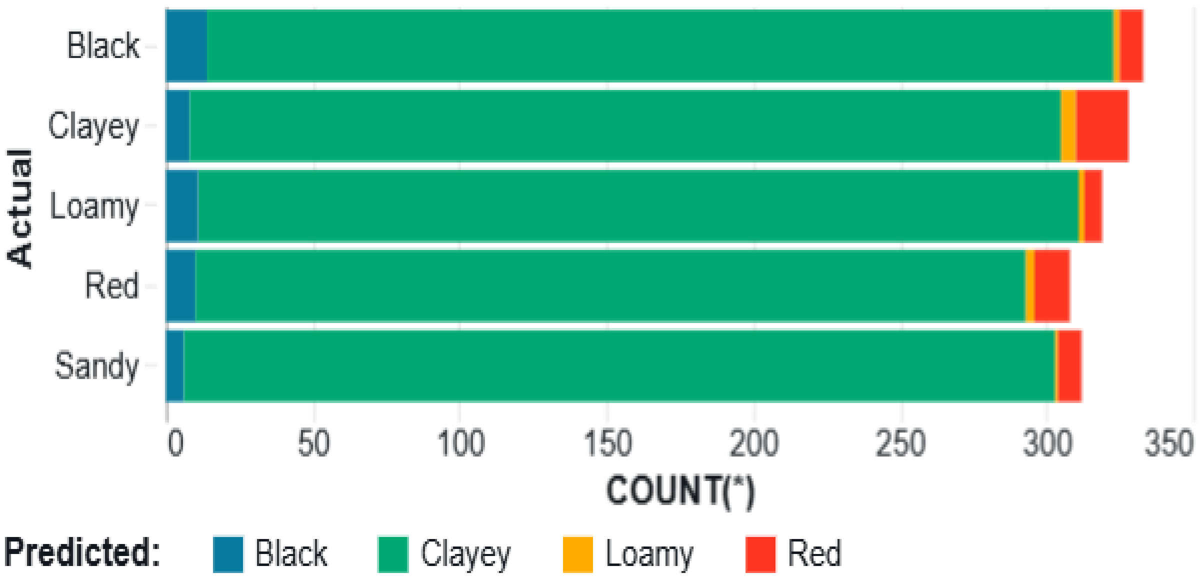

Visualization: Heatmap showing soil–crop relationship and feature importance plot.

Insight: Machine learning–based soil recommendations can support precision agriculture by aligning crop choices with soil health.

Chapter 5 – Conclusion and Future Scope

5.1 Conclusion

This Big Data Capstone Project successfully demonstrated the application of big data analytics, machine learning, and data engineering techniques in the field of agriculture using Databricks and Apache Spark.
The project aimed to provide actionable insights into improving crop productivity, understanding climatic impacts, and recommending suitable crop–soil combinations through five well-defined use cases.

By building a complete data pipeline architecture—from data ingestion (Bronze layer) to transformation (Silver layer), feature engineering, modeling, and visualization—the project showcased how large-scale agricultural data can be processed efficiently and visualized in a single integrated dashboard.

The key outcomes of each use case include:

- **Crop Yield Prediction:** Developed a regression model that accurately predicted crop yields based on rainfall and other climatic variables.
- **Rainfall Impact Analysis:** Identified strong correlations between rainfall levels and agricultural productivity, helping to anticipate drought or flood impacts.
- **Climate Matching for Crops:** Matched different climatic zones with suitable crops, optimizing regional agricultural planning.
- **Growing Degree Days (GDD):** Estimated the ideal crop growth periods to assist farmers in optimizing planting and harvesting schedules.
- **Soil Recommendation Model:** Predicted the best crops for specific soil types, encouraging sustainable and precise farming practices.

Overall, this project proved that big data platforms like Databricks can transform traditional agriculture into data-driven smart farming by integrating multiple data sources and generating accurate, real-time insights.

5.2 Key Achievements

1. Successfully created a **multi-layered pipeline** (Bronze, Silver, Gold) to manage agricultural datasets efficiently.
2. Built **five analytical models** to support yield forecasting, rainfall impact analysis, and soil recommendation.
3. Designed an **interactive Databricks dashboard** for real-time agricultural insights.
4. Improved data quality and performance using **Delta Lake and Spark SQL**.

5. Demonstrated scalability by integrating large datasets in a distributed environment.

5.3 Challenges Faced

1. Handling data inconsistencies and missing values across multiple agricultural datasets.
2. Integrating different data formats (CSV, API-based climate data, soil datasets).
3. Limited data availability for certain regions, affecting model generalization.
4. Complex parameter tuning for machine learning models in PySpark.
5. Ensuring efficient execution of pipelines within Databricks workspace constraints.

5.4 Future Scope

1. **Integration with IoT Devices:** Real-time sensor data (soil moisture, temperature, humidity) can be added to improve model accuracy.
2. **Advanced Machine Learning Models:** Use deep learning or time-series forecasting models like LSTM and ARIMA for yield prediction.
3. **Geospatial Data Integration:** Incorporate satellite imagery and GIS data to analyze crop health and land usage.
4. **Automated Alerts and Recommendations:** Develop a real-time alert system for farmers using APIs that notifies about rainfall, pest risk, or irrigation needs.
5. **Cloud Deployment:** Deploy the pipeline as an end-to-end agricultural analytics service using AWS, Azure, or GCP for wider adoption.
6. **Predictive Decision Support System:** Extend the dashboard into a decision-support web application that assists policymakers and farmers in planning strategies based on live analytics.

5.5 Summary

This project highlights how data analytics and cloud computing can modernize agriculture by integrating datasets, applying predictive analytics, and visualizing outcomes effectively. The Databricks-based pipeline, combined with Spark MLlib and visualization tools, provides a strong foundation for smart agriculture solutions.
With further expansion and real-time data integration, this system can serve as a valuable tool for improving agricultural productivity, sustainability, and resource optimization.